

Task 1 — Install TensorFlow and Keras

Run in terminal / command prompt:

```
python -m pip install --upgrade pip
```

```
pip install tensorflow # installs TF (includes Keras API)
```

Notes:

tensorflow package includes Keras as tf.keras.

If you have a GPU and want GPU support, install proper CUDA/cuDNN and use a compatible TF version (check TensorFlow docs for matching versions).

Task 2 — Perceptron model for binary classification (simple example)

This trains a single-layer perceptron (logistic regression style) on a synthetic binary dataset.

```
# perceptron_binary.py
```

```
import numpy as np
```

```
from sklearn.datasets import make_classification
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
# 1) Create synthetic binary classification data
```

```
X, y = make_classification(n_samples=2000, n_features=10, n_informative=6,  
                           n_redundant=2, n_clusters_per_class=2, random_state=42)
```

```
# 2) Train-test split and scaling
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
# 3) Build simple perceptron model
```

```
model = models.Sequential([  
    layers.Input(shape=(X_train.shape[1],)),  
    layers.Dense(1, activation='sigmoid') # single neuron -> perceptron (logistic)  
])
```

```
model.compile(optimizer='adam',
```

```
               loss='binary_crossentropy',
```

```
               metrics=['accuracy'])
```

```
# 4) Train
```

```
history = model.fit(X_train, y_train, epochs=30, batch_size=32, validation_split=0.1,
verbose=2)
```

```
# 5) Evaluate
```

```
loss, acc = model.evaluate(X_test, y_test, verbose=0)
```

```
print(f"Test loss: {loss:.4f}, Test accuracy: {acc:.4f}")
```

Explanation: A perceptron with sigmoid outputs probability for class 1; trained with binary cross-entropy.

Task 3 — Multi-Layer Perceptron (MLP) for MNIST

Simple MLP that classifies MNIST handwritten digits using `tf.keras.datasets`.

```
# mlp_mnist.py
```

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
# 1) Load data
```

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

```
# 2) Preprocess: normalize and flatten
```

```
x_train = x_train.astype('float32') / 255.0
```

```
x_test = x_test.astype('float32') / 255.0
```

```
x_train = x_train.reshape((-1, 28*28))
```

```
x_test = x_test.reshape((-1, 28*28))
```

```
# 3) Build MLP
```

```
model = models.Sequential([  
    layers.Input(shape=(28*28,)),  
    layers.Dense(128, activation='relu'),  
    layers.Dense(64, activation='relu'),  
    layers.Dense(10, activation='softmax')  
])
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# 4) Train
```

```
history = model.fit(x_train, y_train, epochs=10, batch_size=128, validation_split=0.1,  
                    verbose=2)
```

```
# 5) Evaluate
```

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)  
print(f"MNIST test acc: {test_acc:.4f}")
```

Tip: You can increase epochs or use dropout/regularization to improve generalization.

Task 4 — Visualize model accuracy and loss graphs using Matplotlib

Use the history object returned by `model.fit()` to plot.

visualize_history.py (use after training to plot the history returned by fit)

```
import matplotlib.pyplot as plt
```

```
def plot_history(history):
```

```
    # loss
```

```
    plt.figure(figsize=(10,4))
```

```
    plt.subplot(1,2,1)
```

```
    plt.plot(history.history['loss'], label='train_loss')
```

```
    plt.plot(history.history.get('val_loss', []), label='val_loss')
```

```
    plt.xlabel('Epoch')
```

```
    plt.ylabel('Loss')
```

```
    plt.legend()
```

```
    plt.title('Loss vs Epoch')
```

```
    # accuracy
```

```
    plt.subplot(1,2,2)
```

```
    plt.plot(history.history.get('accuracy', history.history.get('acc')), label='train_acc')
```

```
    plt.plot(history.history.get('val_accuracy', history.history.get('val_acc', [])), label='val_acc')
```

```
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy vs Epoch')
```

```
plt.tight_layout()
plt.show()
```

```
# Example usage (after training): plot_history(history)
```

If using this in a notebook, `plt.show()` will render plots inline.

Task 5 — Perform hyperparameter tuning (change epochs, batch size, activation functions)

Simple manual grid search example (train small model multiple times and compare). For quick experiments, use a small subset or fewer epochs. For production-scale tuning, use `keras-tuner`.

```
# simple_hyperparam_search.py

import itertools

import tensorflow as tf

from tensorflow.keras import layers, models

from tensorflow.keras.datasets import mnist

import numpy as np
```

```
# use a small subset to speed up grid search in examples
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
x_train = x_train.astype('float32') / 255.0
```

```
x_test = x_test.astype('float32') / 255.0
```

```
x_train = x_train.reshape((-1, 28*28))
```

```
x_test = x_test.reshape((-1, 28*28))
```

```
# Use a small subset for quick experiments
```

```
x_train_small = x_train[:20000]
```

```
y_train_small = y_train[:20000]
```

```
x_val = x_train[20000:25000]
```

```
y_val = y_train[20000:25000]
```

```
def build_model(activation='relu'):
```

```
    model = models.Sequential([
```

```
        layers.Input(shape=(28*28,)),
```

```
        layers.Dense(128, activation=activation),
```

```
        layers.Dense(64, activation=activation),
```

```
        layers.Dense(10, activation='softmax')
```

```
    ])
```

```
    model.compile(optimizer='adam',
```

```
                  loss='sparse_categorical_crossentropy',
```

```
                  metrics=['accuracy'])
```

```
    return model
```

```

# hyperparameters to try
param_grid = {
    'epochs': [5, 10],
    'batch_size': [64, 128],
    'activation': ['relu', 'tanh']
}

results = []

for epochs, batch_size, activation in itertools.product(param_grid['epochs'],
                                                         param_grid['batch_size'],
                                                         param_grid['activation']):

    print(f"Training: epochs={epochs}, batch_size={batch_size}, activation={activation}")
    model = build_model(activation=activation)
    hist = model.fit(x_train_small, y_train_small, epochs=epochs, batch_size=batch_size,
                    validation_data=(x_val, y_val), verbose=0)
    val_acc = hist.history['val_accuracy'][-1]
    results.append({
        'epochs': epochs, 'batch_size': batch_size, 'activation': activation,
        'val_acc': val_acc
    })

    print(f" -> val_acc={val_acc:.4f}")

# show best
best = max(results, key=lambda r: r['val_acc'])
print("Best config:", best)

```


Alternative: Use keras-tuner for a more principled search (random search / Bayesian optimization).

Task 6 — Short blog/note explaining how neural networks learn

English (short):

Neural networks learn by adjusting weights to reduce the difference between predicted outputs and target outputs. During training:

1. A forward pass computes predictions from inputs through layers (matrices multiply and nonlinear activations).
2. A loss (error) function quantifies how far predictions are from true labels (e.g., cross-entropy).
3. Backpropagation computes gradients of the loss w.r.t. each weight using the chain rule.
4. An optimizer (e.g., SGD, Adam) updates weights in the direction that reduces loss using those gradients.

5. Repeat over many examples (epochs) until the network generalizes well.

Key concepts:

Activation functions (ReLU, sigmoid, tanh) add nonlinearity so networks can learn complex mappings.

Learning rate controls step size of updates; too large → divergence; too small → slow training.

Overfitting vs generalization: networks can memorize training data — regularization (dropout, weight decay), early stopping, and more data help generalize.

Batch size affects gradient estimate quality: small batches are noisy but can generalize well; large batches give stable gradients.